App Deployment Inside VPC On Private Server in Production

About the Project:

This example demonstrates how to create a VPC that you use for servers in a production environment.

To improve resiliency, you deploy the servers in two AZs by using an auto scaling group and an application load balancer. For additional security, you deploy the servers in private subnets. The server receives requests through the load balancer. The servers can connect to the internet gateway by using a NAT gateway. To improve resiliency, you deploy the NAT gateway in both availability zones.

Project Overview:

The VPC has public subnets and private subnets in two availability zones.

Each public subnet contains a NAT gateway and a load balancer.

The servers run in the private subnets are launched and terminated by using an auto-scaling group, and receives traffic from the load balancer.

The server can connect to the internet by using NAT gateway.

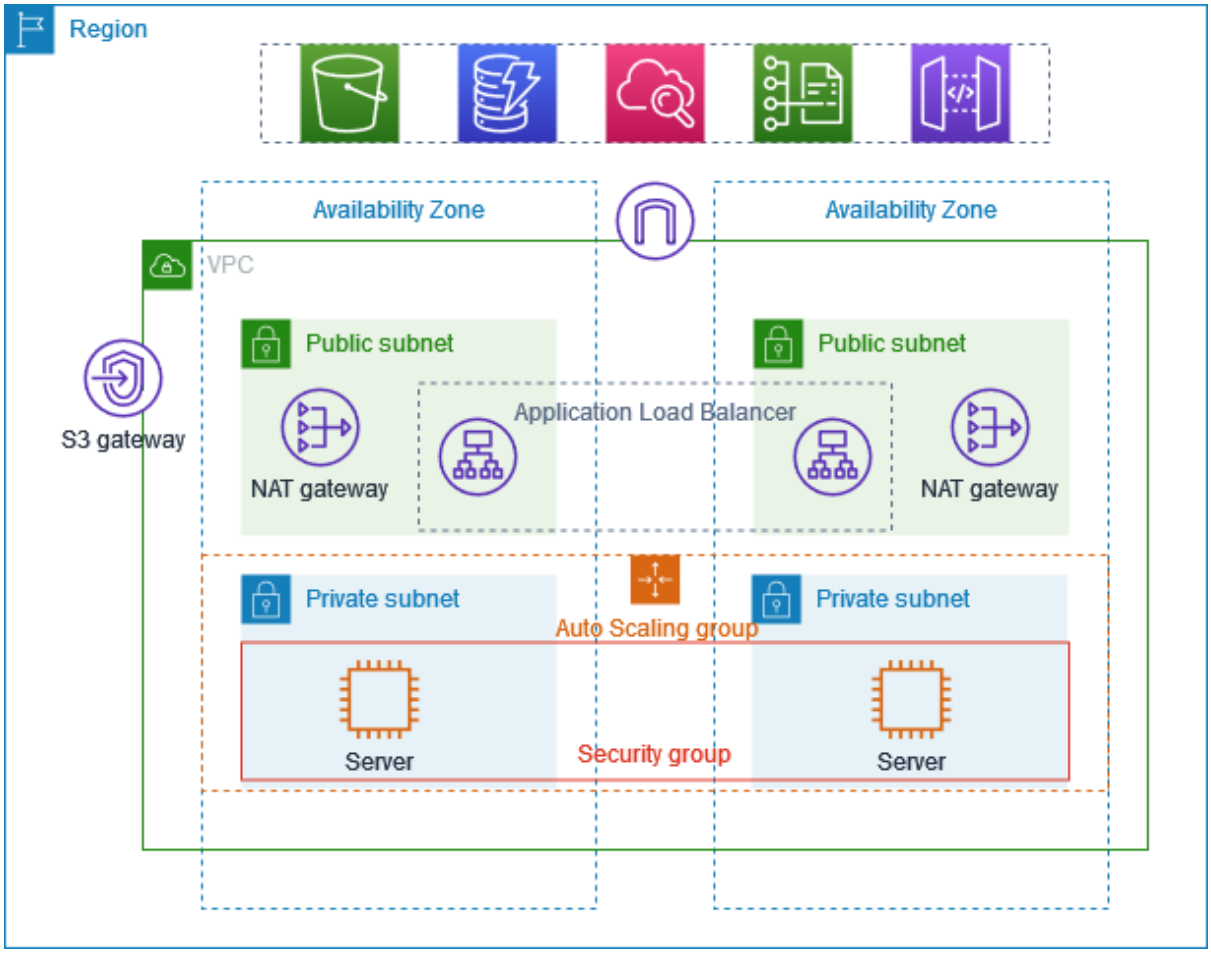

1. Go to VPC service and click on create VPC.
2. Follow the configuration and create a VPC as per your need.

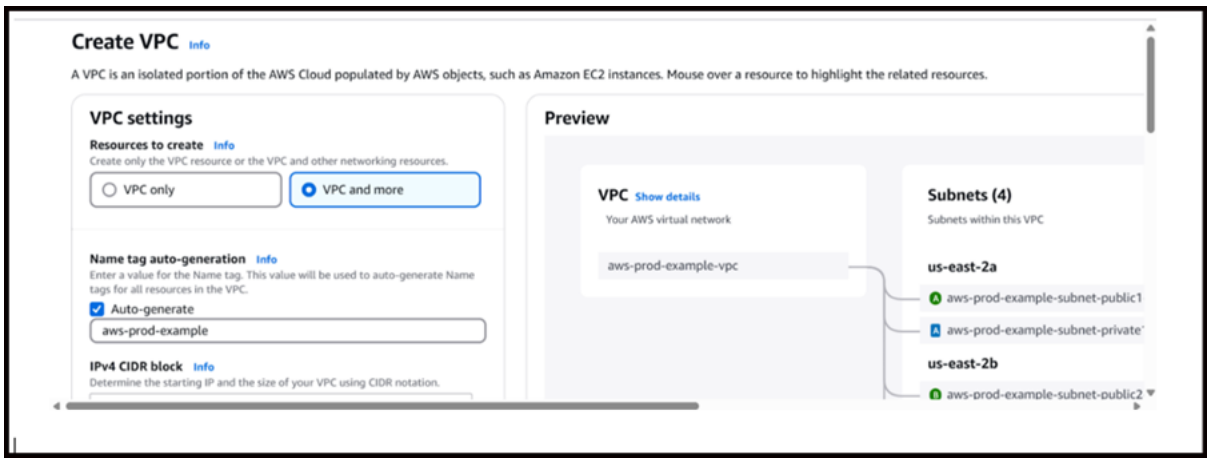

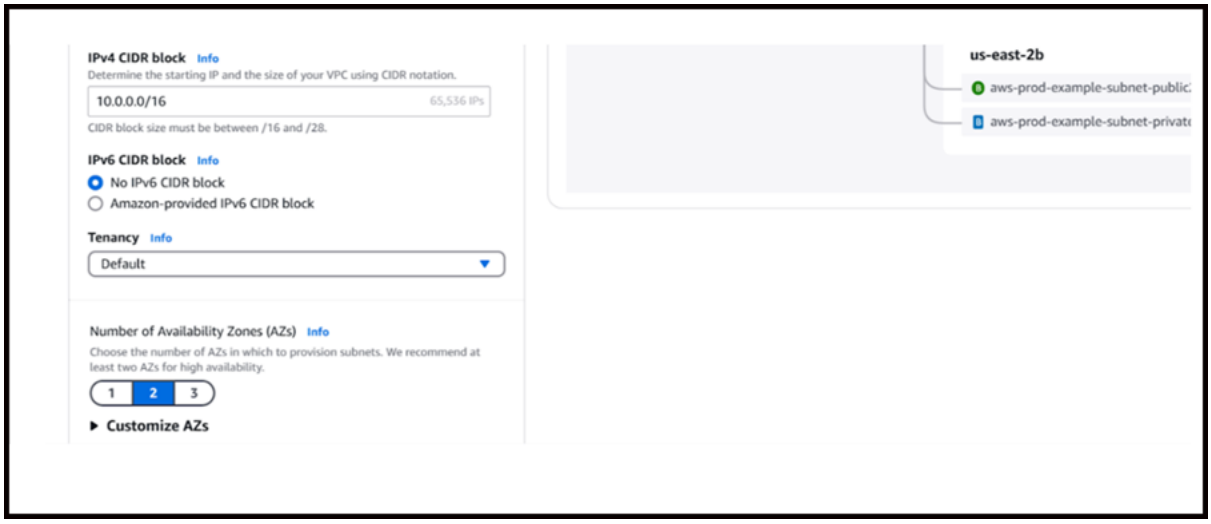

Now click on create VPC.

This would be the architecture of our VPC

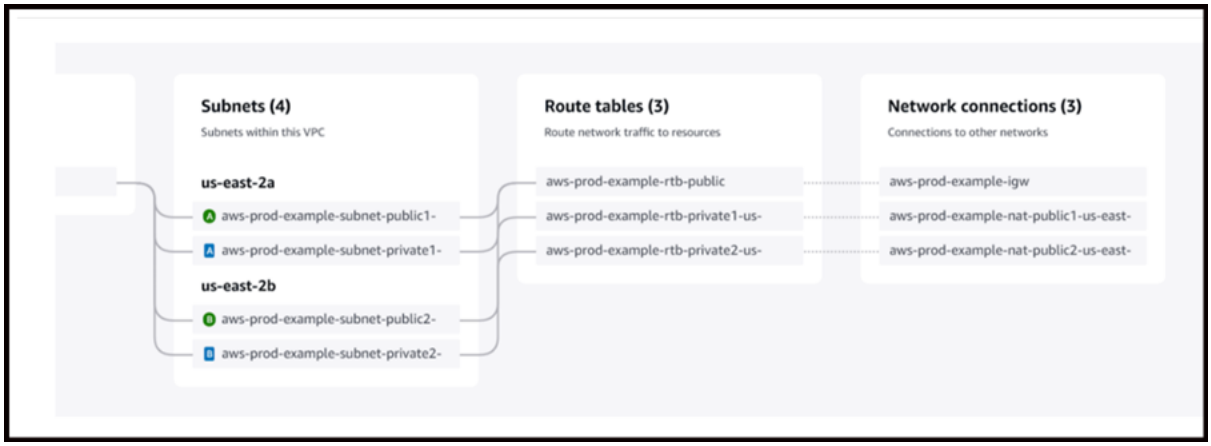

3. Now create two EC2 instances with auto scaling and a load balancer. Go to the auto scaling group and create an auto-scaling group.

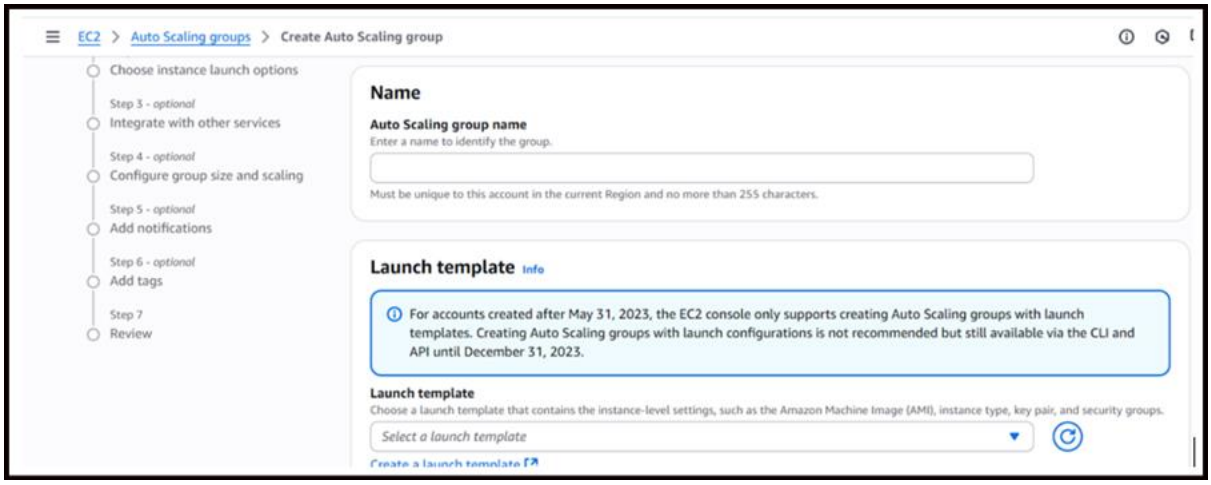

4. Click on create a launch templates and select following configurations:

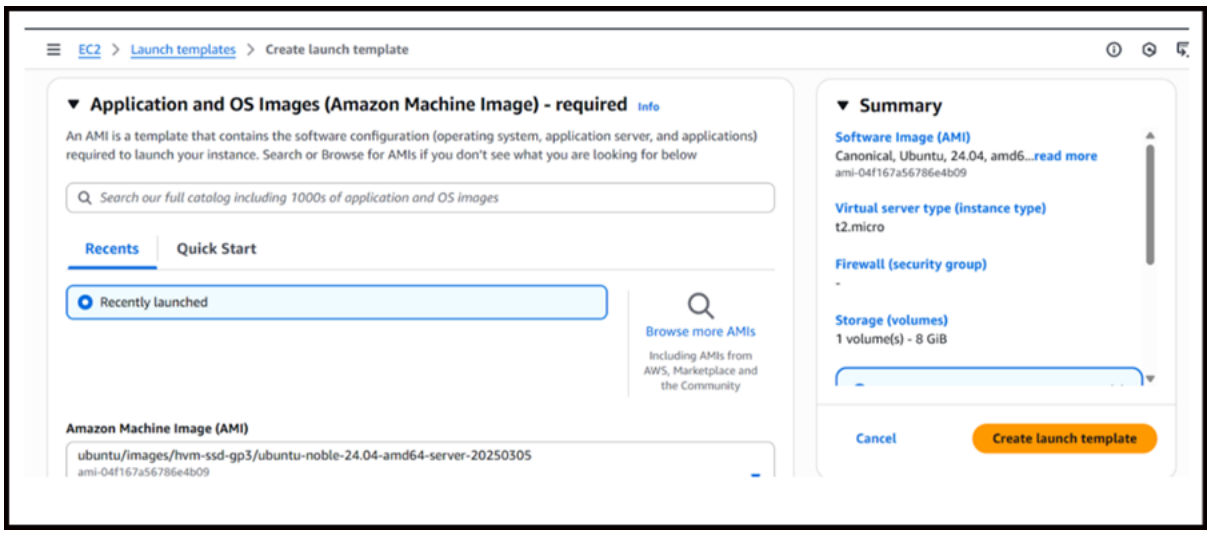

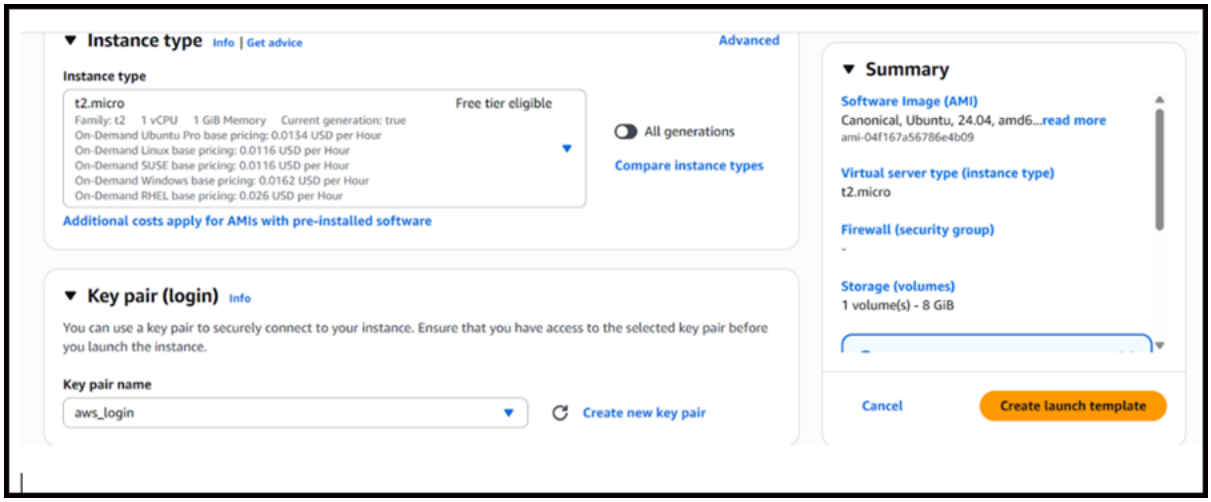
Instance type Info | Get advice
Advanced
Instance type
t2.micro
Free tier eligible
Family: t2 1 vCPU 1 GiB Memory Current generation: true
On-Demand Ubuntu Pro base pricing: 0.0134 USD per Hour
On-Demand Linux base pricing: 0.0116 USD per Hour
On-Demand SUSE base pricing: 0.0116 USD per Hour
On-Demand Windows base pricing: 0.0162 USD per Hour
On-Demand RHEL base pricing: 0.026 USD per Hour
All generations
Compare instance types
Additional costs apply for AMIs with pre-installed software
Key pair (login) Info
You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.
Key pair name
aws_login
Create new key pair
Summary
Software Image (AMI)
Canonical, Ubuntu, 24.04, amd6...read more
ami-04f167a56786e4b09
Virtual server type (instance type)
t2.micro
Firewall (security group)
-
Storage (volumes)
1 volume(s) - 8 GiB
Cancel
Create launch template

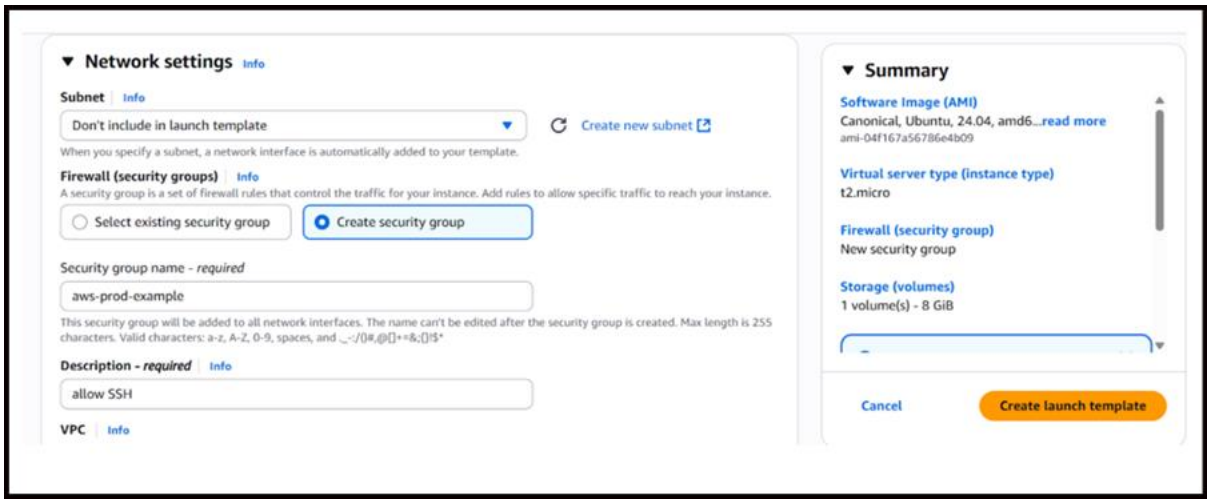
Network settings Info
Subnet Info
Don't include in launch template
Create new subnet
When you specify a subnet, a network interface is automatically added to your template.
Firewall (security groups) Info
A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.
Select existing security group
Create security group
Security group name - required
aws-prod-example
This security group will be added to all network interfaces. The name can't be edited after the security group is created. Max length is 255 characters. Valid characters: a-z, A-Z, 0-9, spaces, and ._-:/()#,@[]+=&;{}!$*
Description - required Info
allow SSH
VPC Info
Summary
Software Image (AMI)
Canonical, Ubuntu, 24.04, amd6...read more
ami-04f167a56786e4b09
Virtual server type (instance type)
t2.micro
Firewall (security group)
New security group
Storage (volumes)
1 volume(s) - 8 GiB
Cancel
Create launch template

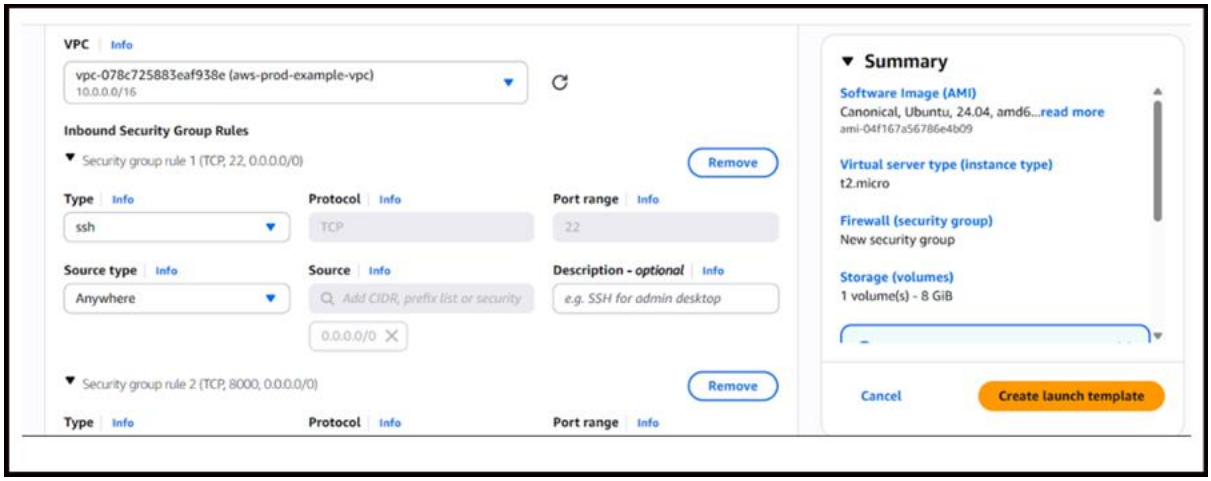
VPC Info
vpc-078c725883eaf938e (aws-prod-example-vpc)
10.0.0.0/16
Inbound Security Group Rules
Security group rule 1 (TCP, 22, 0.0.0.0/0)
Remove
Type Info
ssh
Protocol Info
TCP
Port range Info
22
Source type Info
Anywhere
Source Info
Add CIDR, prefix list or security
0.0.0.0/0
Description - optional Info
e.g. SSH for admin desktop
Security group rule 2 (TCP, 8000, 0.0.0.0/0)
Remove
Type Info
Protocol Info
Port range Info
Summary
Software Image (AMI)
Canonical, Ubuntu, 24.04, amd6...read more
ami-04f167a56786e4b09
Virtual server type (instance type)
t2.micro
Firewall (security group)
New security group
Storage (volumes)
1 volume(s) - 8 GiB
Cancel
Create launch template

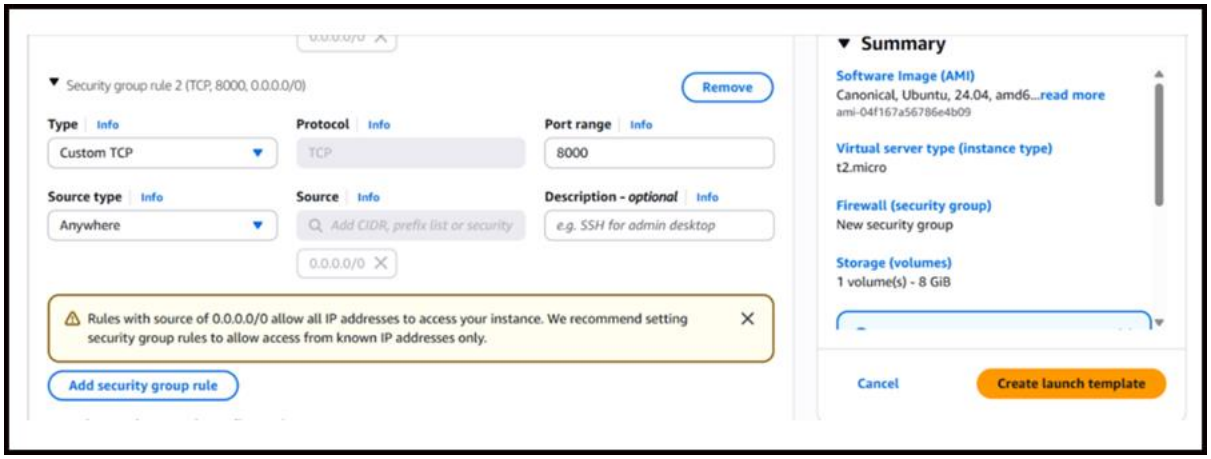

And click on the launch template.

5. Do refresh the page and go back to the auto-scaling group and fill the necessary thing.

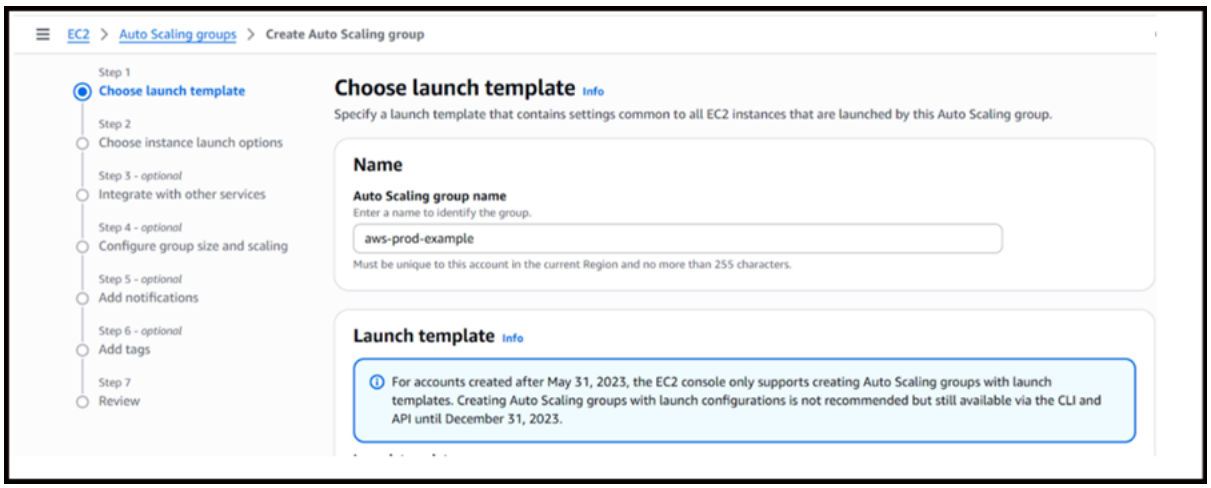

and choose the template you created and click on the next button.

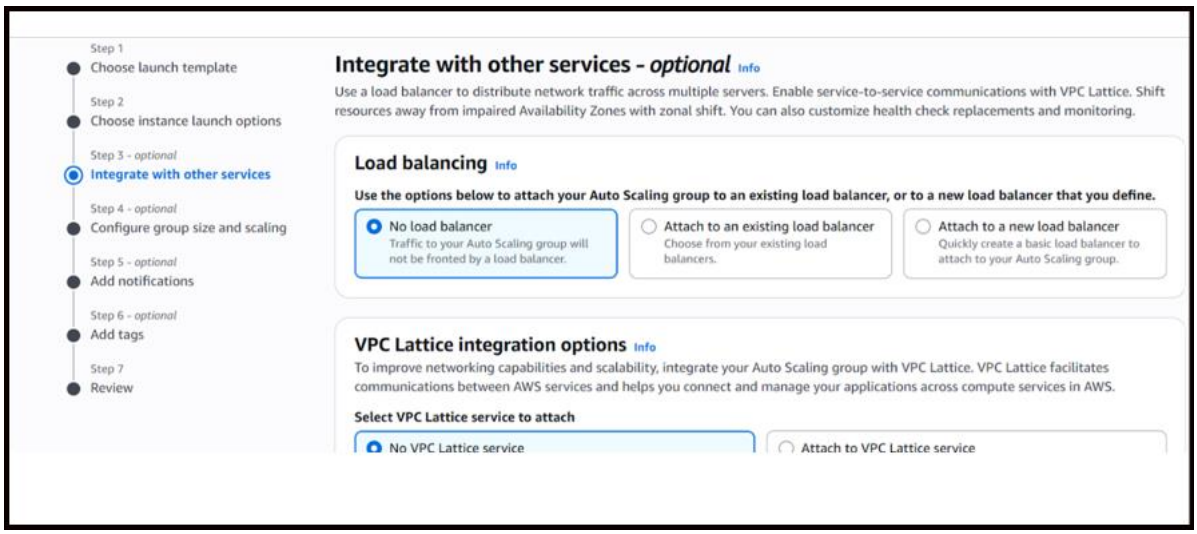

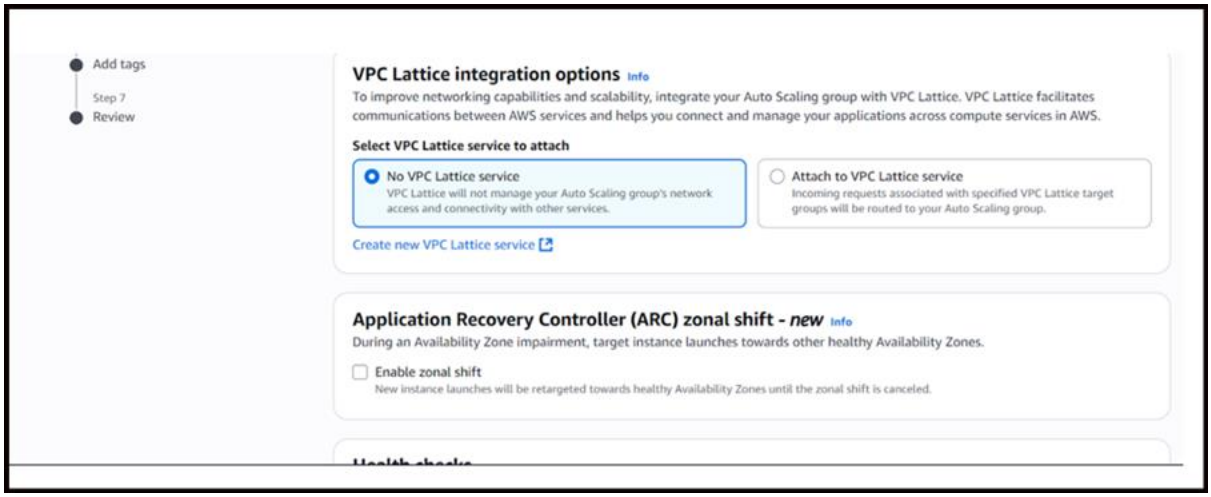
Add tags
Step 7
Review
VPC Lattice integration options Info
To improve networking capabilities and scalability, integrate your Auto Scaling group with VPC Lattice. VPC Lattice facilitates communications between AWS services and helps you connect and manage your applications across compute services in AWS.
Select VPC Lattice service to attach
No VPC Lattice service
VPC Lattice will not manage your Auto Scaling group's network access and connectivity with other services.
Attach to VPC Lattice service
Incoming requests associated with specified VPC Lattice target groups will be routed to your Auto Scaling group.
Create new VPC Lattice service
Application Recovery Controller (ARC) zonal shift - new Info
During an Availability Zone impairment, target instance launches towards other healthy Availability Zones.
Enable zonal shift
New instance launches will be retargeted towards healthy Availability Zones until the zonal shift is canceled.

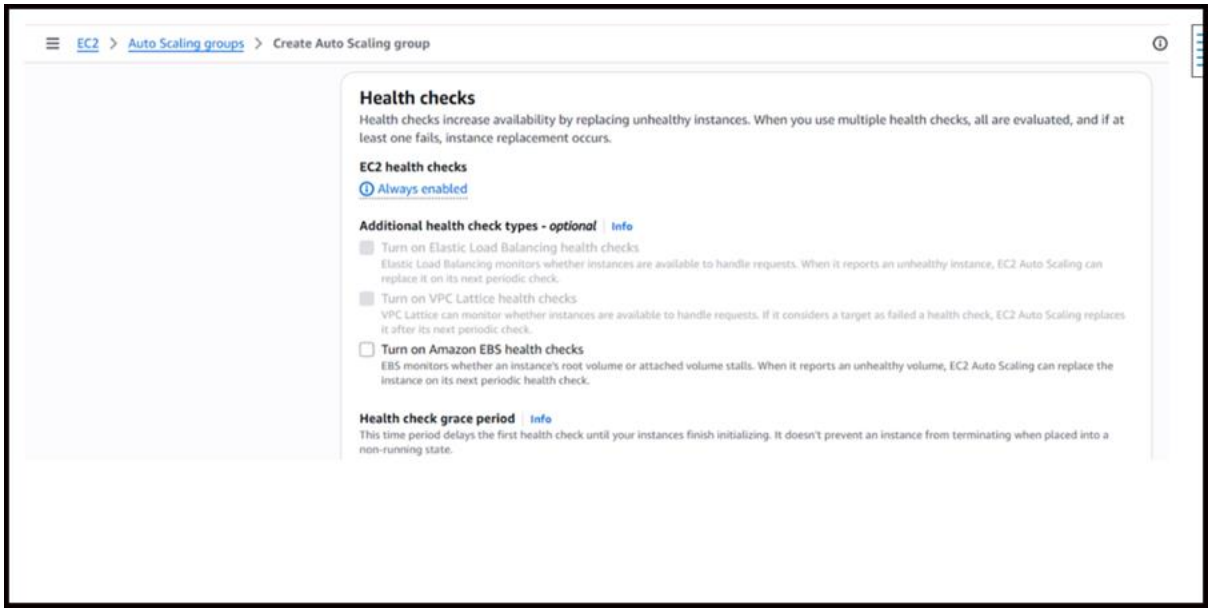
EC2 > Auto Scaling groups > Create Auto Scaling group
Health checks
Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.
EC2 health checks
Always enabled
Additional health check types - optional Info
Turn on Elastic Load Balancing health checks
Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.
Turn on VPC Lattice health checks
VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.
Turn on Amazon EBS health checks
EBS monitors whether an instance's root volume or attached volume stalls. When it reports an unhealthy volume, EC2 Auto Scaling can replace the instance on its next periodic health check.
Health check grace period Info
This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.

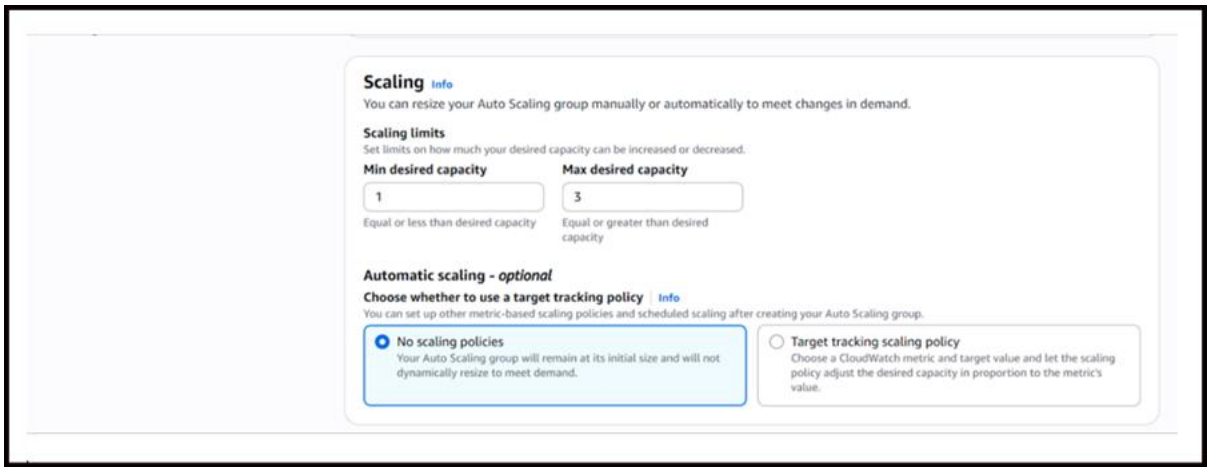
Scaling Info
You can resize your Auto Scaling group manually or automatically to meet changes in demand.
Scaling limits
Set limits on how much your desired capacity can be increased or decreased.
Min desired capacity
1
Equal or less than desired capacity
Max desired capacity
3
Equal or greater than desired capacity
Automatic scaling - optional
Choose whether to use a target tracking policy Info
You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.
No scaling policies
Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.
Target tracking scaling policy
Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

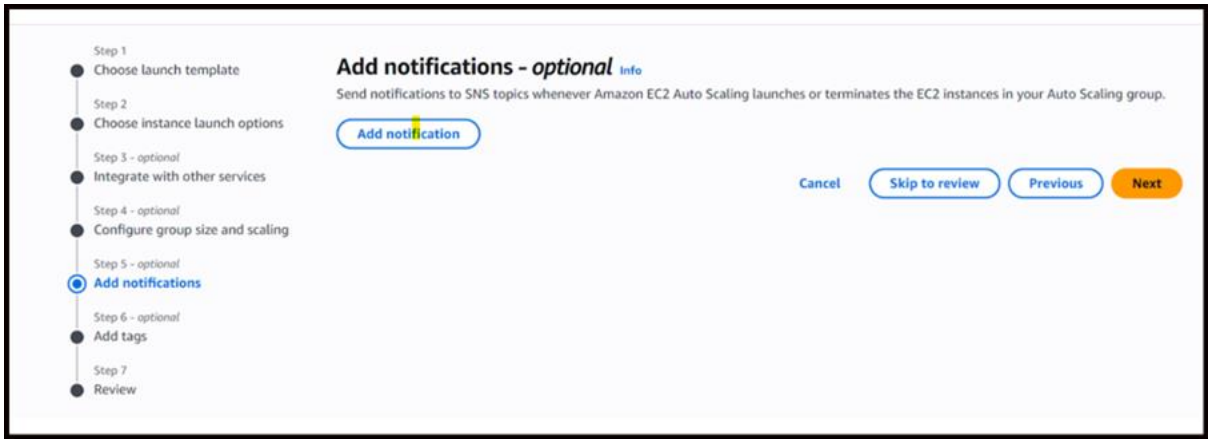

click on the next button after doing all above configurations. Verify all your configurations and click on create auto-scaling group.

Two instances got created in us-east-2a and us-east-2b

These instances don't have an IP address in public because we created a private subnets server. So we use a jump server approach to connect to them.

Jump Server - In AWS, a jump server (also called a bastion host) is a dedicated server that acts as a secure gateway, providing access to private instances within a VPC (Virtual Private Cloud) from outside the VPC. It's a common security practice that reduces the risk of unauthorized access to internal resources by funneling all external SSH traffic through a single, secure point.

For this go to launch instance and create a new instance as bastion host or jump server with created VPC and public subnets and give access to SSH login.

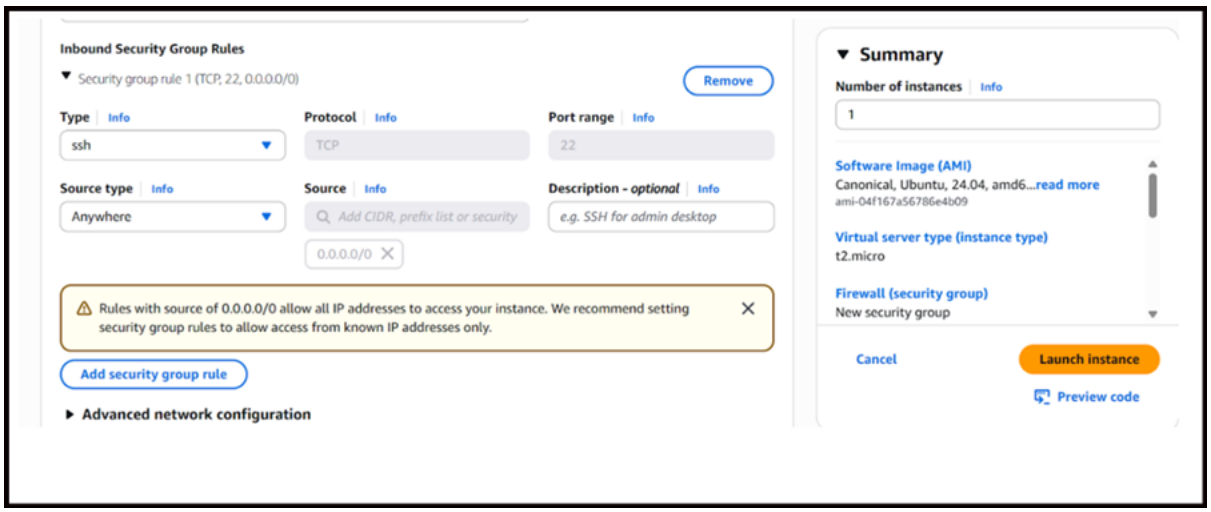

And launch the instance. Now you will have 3 instances.

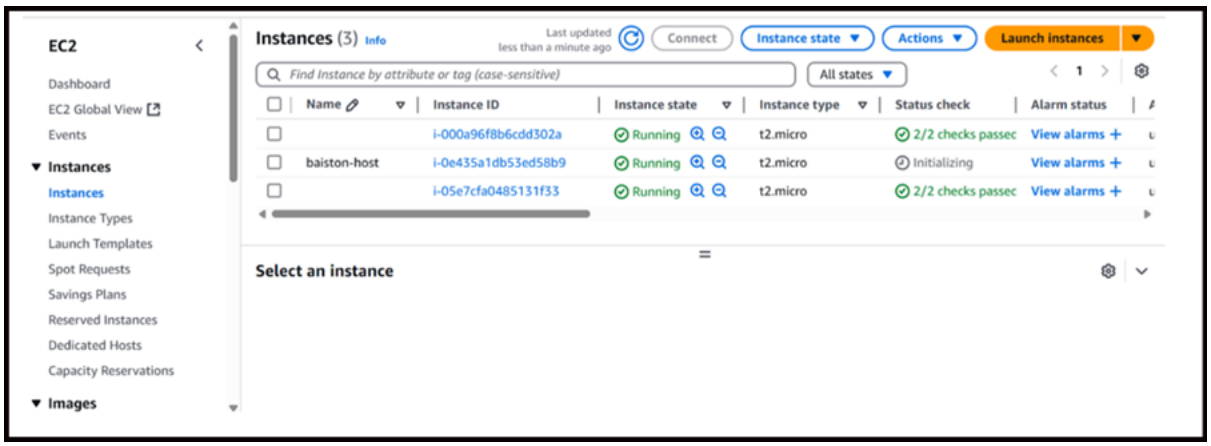

2 created with an auto-scaling group and one as a jump server.

To connect with private subnets, your jump server should connect to the SSH because we want to connect 2 instances created by auto-scaling.

So we have to copy the pem file of 2 instances from our local server to the jump server.SCP Command - scp -i "aws_login.pem" aws_login.pem ubuntu@ec2-18-116-80-56.us-east-2.compute.amazonaws.com:/home/ubuntu

After copying file login to the jump server by using SSH key.

This pem file should be here in your jump server.

Now go to any one instance and take the private IP address and launch an application.

And do ssh in instance.

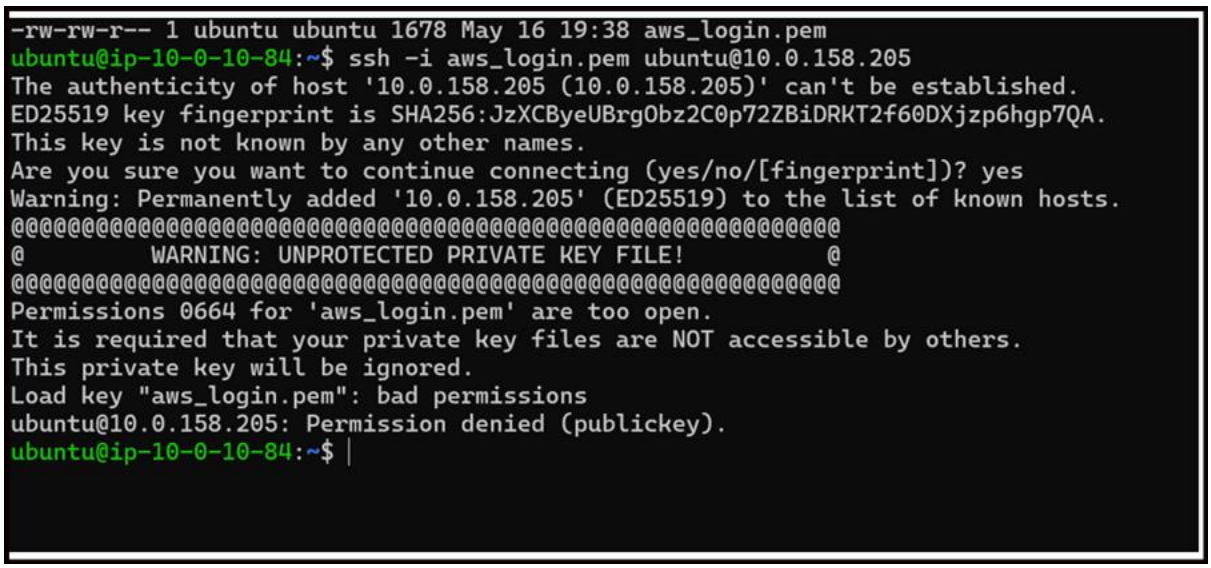

Change the permission of pem to login to the server.

chmod 600 aws_login.pem

and do ssh -i aws_login.pem ubuntu@10.0.158.205 and get connected to the server.

```
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/pro

 System information as of Fri May 16 19:46:35 UTC 2025

  System load:  0.0               Processes:             103
  Usage of /:   25.0% of 6.71GB   Users logged in:       0
  Memory usage: 20%               IPv4 address for enX0: 10.0.158.205
  Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status


The list of available updates is more than a week old.
To check for new updates run: sudo apt update


The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-10-0-158-205:~$
```

create an index.html and write a basic html inside the index.html file.

```
ubuntu@ip-10-0-158-205:~$ ls -l
total 0
ubuntu@ip-10-0-158-205:~$ vim index.html
```

```
ubuntu@ip-10-0-158-205:~$ ls -l
total 0
ubuntu@ip-10-0-158-205:~$ vim index.html
ubuntu@ip-10-0-158-205:~$ vim index.html
ubuntu@ip-10-0-158-205:~$ python3 --version
Python 3.12.3
ubuntu@ip-10-0-158-205:~$ python3 -m http.server 8000
Serving HTTP on 0.0.0.0 port 8000 (http://0.0.0.0:8000/) ...
```

Now go to the load balancers section and create a default application load balancer which moves all your traffic to application (python).

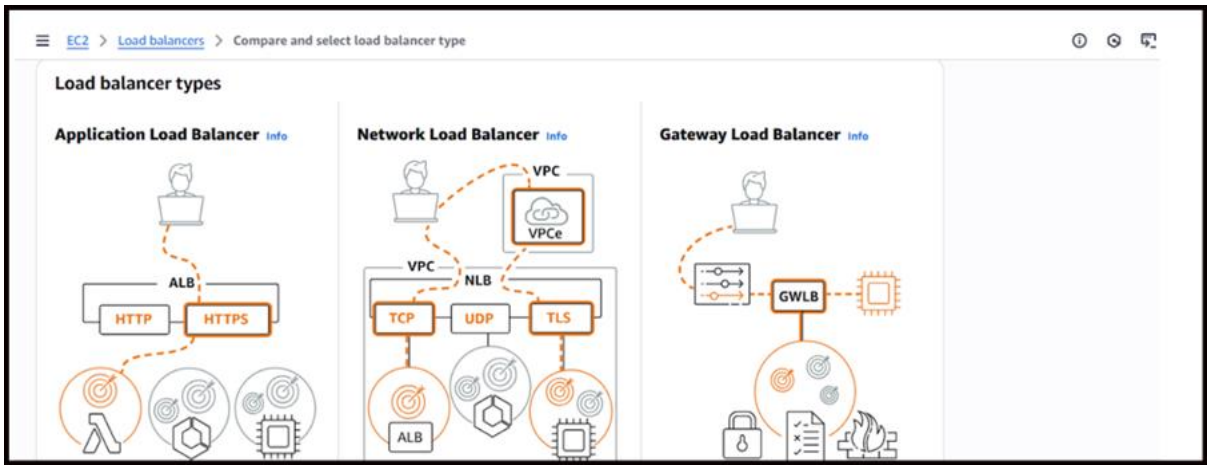

Application load balancer - The Application Load Balancer distributes incoming HTTP and HTTPS traffic across multiple targets such as Amazon EC2 instances, microservices, and containers, based on request attributes. When the load balancer receives a connection request, it evaluates the listener rules in priority order to determine which rule to apply, and if applicable, it selects a target from the target group for the rule action.

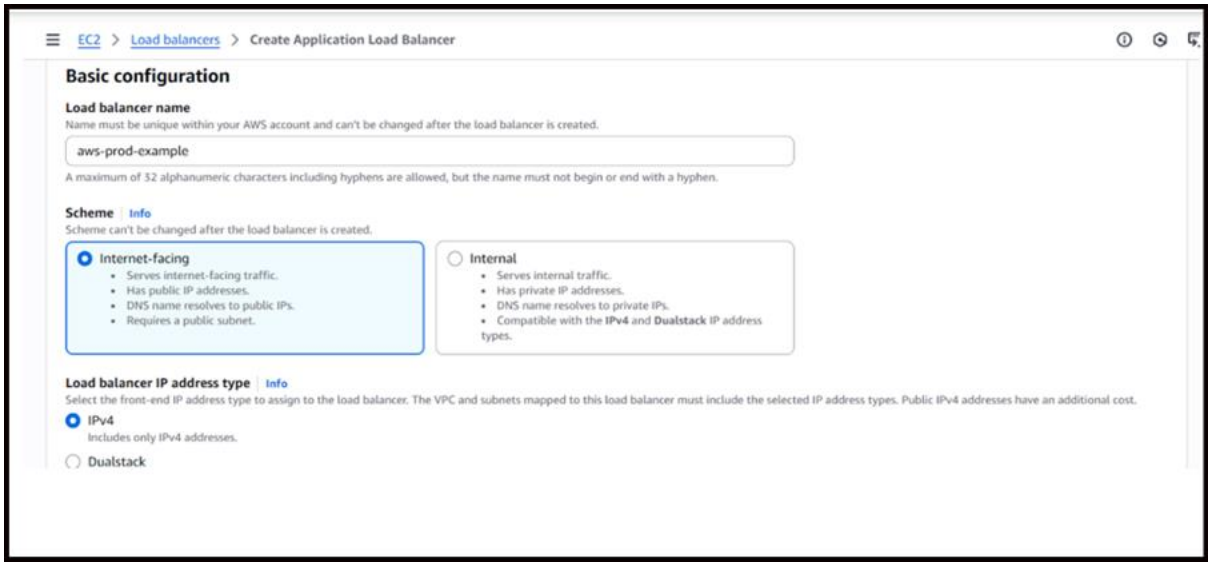

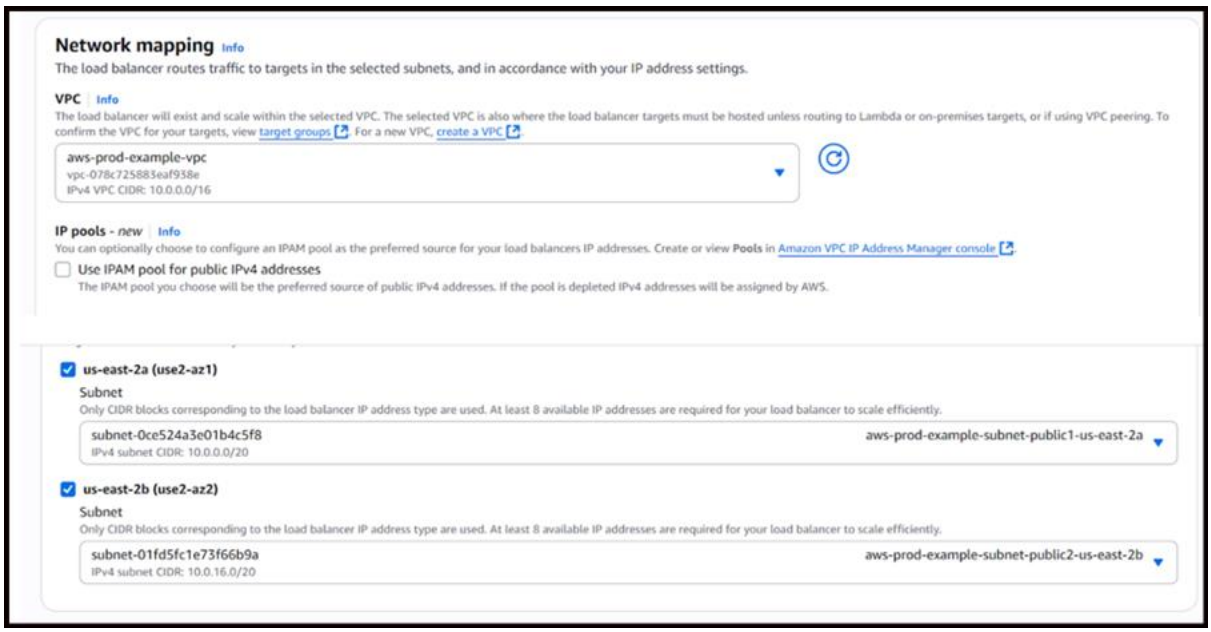

Give the public subnet in subnet section because in our diagram load balancer attached to the public subnet.

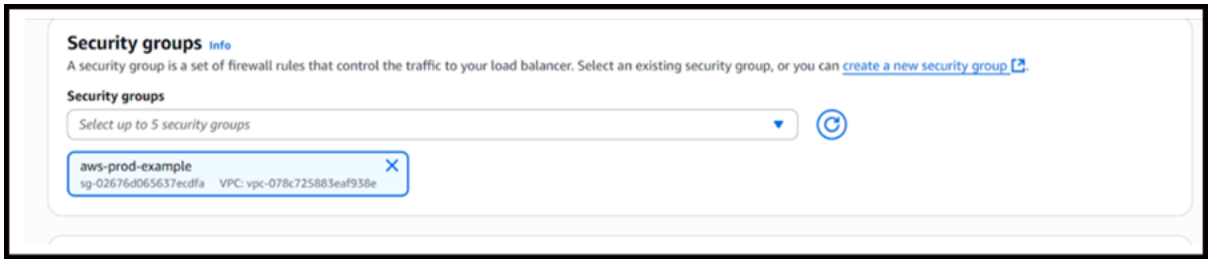

create a target group for listening to port 8000 and after configuring all the things click on next.

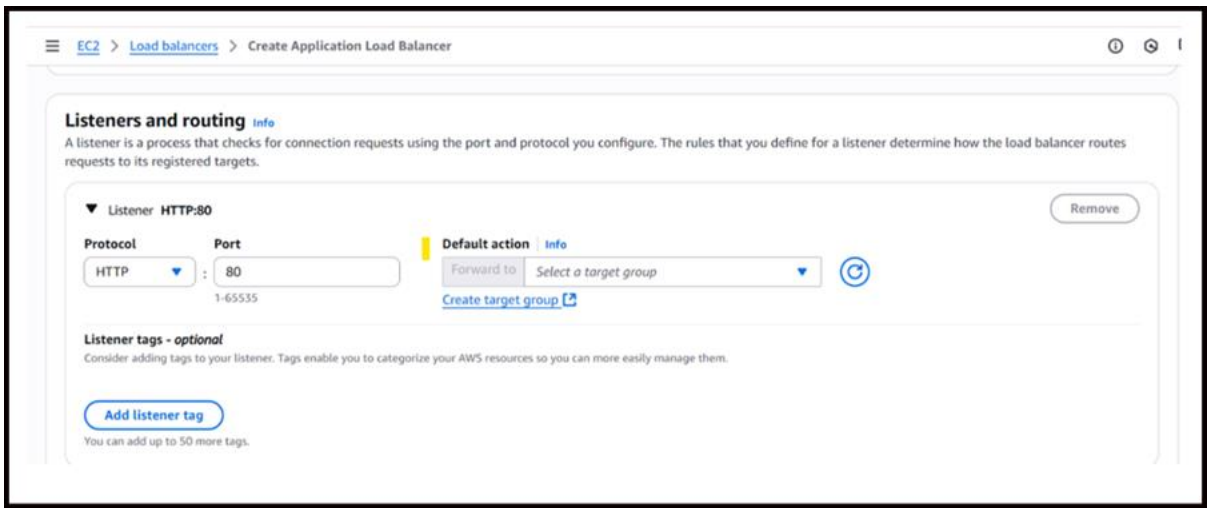
EC2 > Load balancers > Create Application Load Balancer
Listeners and routing Info
A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.
Listener HTTP:80
Remove
Protocol
Port
HTTP
80
1-65535
Default action Info
Forward to
Select a target group
Create target group
Listener tags - optional
Consider adding tags to your listener. Tags enable you to categorize your AWS resources so you can more easily manage them.
Add listener tag
You can add up to 50 more tags.

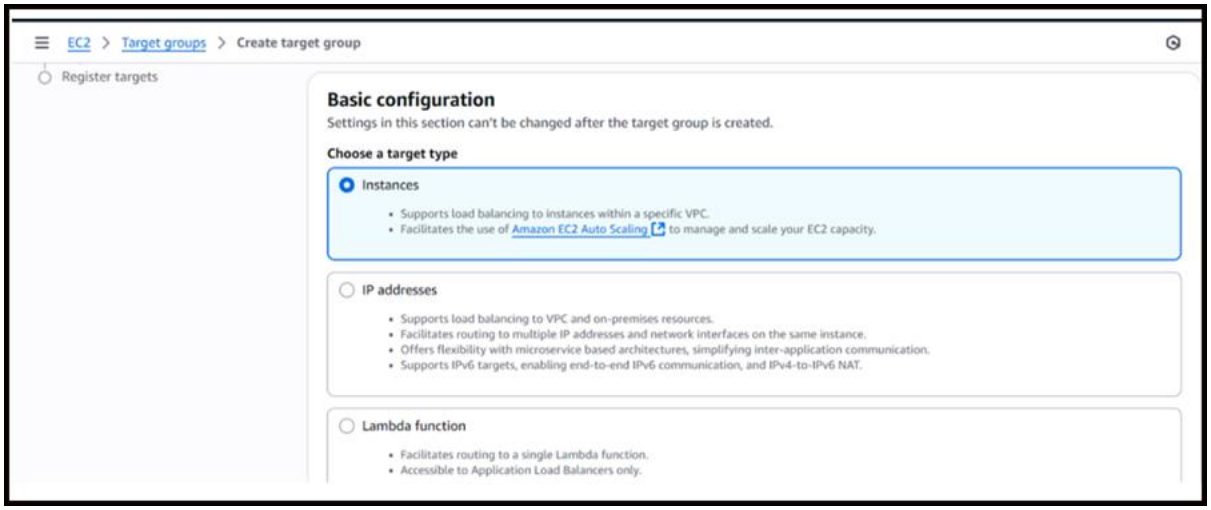
EC2 > Target groups > Create target group
Register targets
Basic configuration
Settings in this section can't be changed after the target group is created.
Choose a target type
Instances
Supports load balancing to instances within a specific VPC.
Facilitates the use of Amazon EC2 Auto Scaling to manage and scale your EC2 capacity.
IP addresses
Supports load balancing to VPC and on-premises resources.
Facilitates routing to multiple IP addresses and network interfaces on the same instance.
Offers flexibility with microservice based architectures, simplifying inter-application communication.
Supports IPv6 targets, enabling end-to-end IPv6 communication, and IPv4-to-IPv6 NAT.
Lambda function
Facilitates routing to a single Lambda function.
Accessible to Application Load Balancers only.

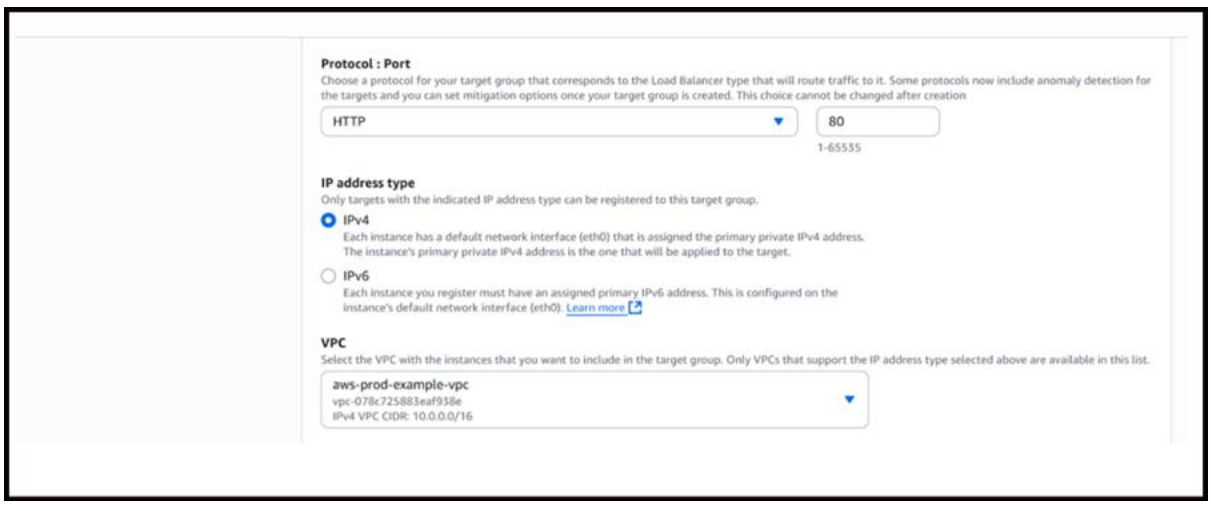
Protocol : Port
Choose a protocol for your target group that corresponds to the Load Balancer type that will route traffic to it. Some protocols now include anomaly detection for the targets and you can set mitigation options once your target group is created. This choice cannot be changed after creation
HTTP
80
1-65535
IP address type
Only targets with the indicated IP address type can be registered to this target group.
IPv4
Each instance has a default network interface (eth0) that is assigned the primary private IPv4 address.
The instance's primary private IPv4 address is the one that will be applied to the target.
IPv6
Each instance you register must have an assigned primary IPv6 address. This is configured on the instance's default network interface (eth0). Learn more
VPC
Select the VPC with the instances that you want to include in the target group. Only VPCs that support the IP address type selected above are available in this list.
aws-prod-example-vpc
vpc-078c725883eaf938e
IPv4 VPC CIDR: 10.0.0.0/16

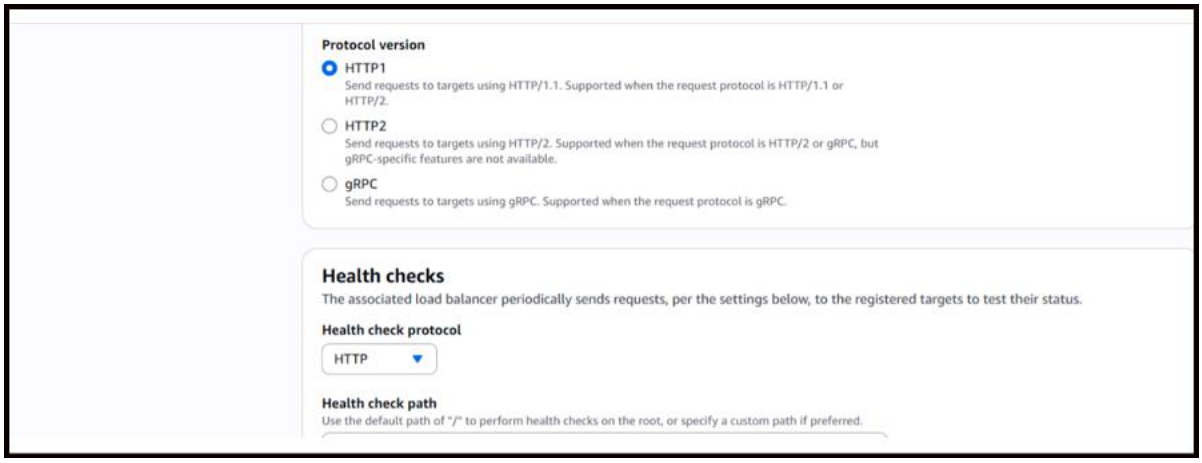

It will open a instance section and choose both the instance because one has an application and other not we have to check the incoming traffic of private subnets in one server and other one throw an error because it doesn't have the application.

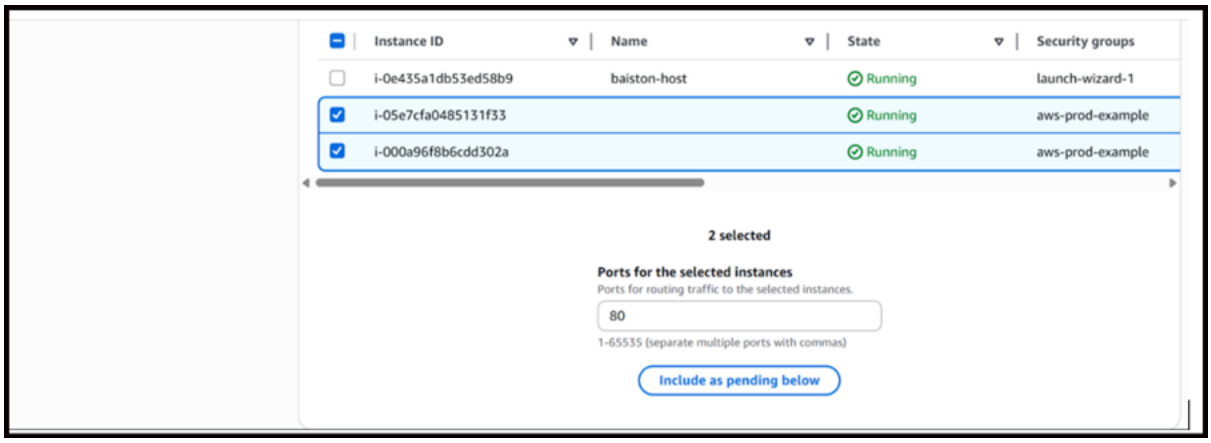

Include as pending below, click on that and it shows the target instance where the flow of incoming traffic should go and create target groups.

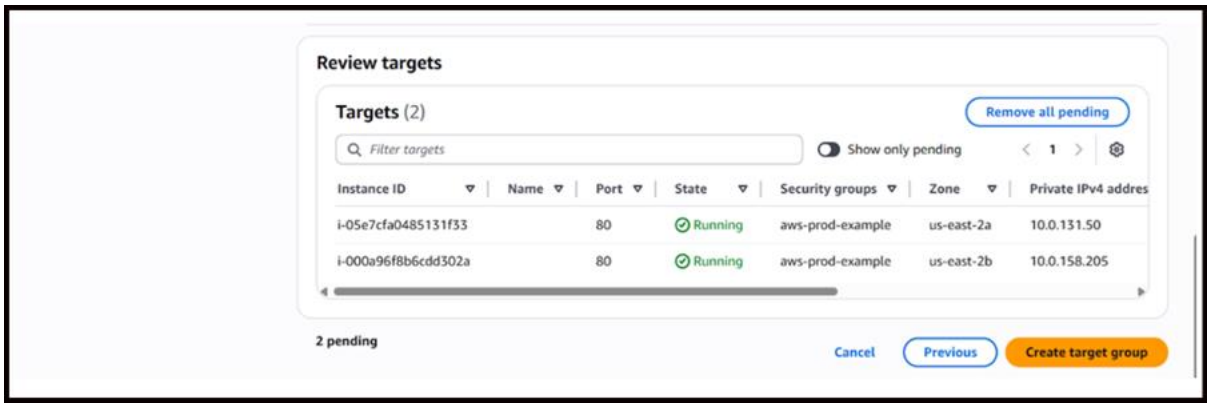

Run your application on port instead of 8000. Relaunch your application in port 80 likewise you can configure 8000 also.

Now go back to the load balancer and do configure for 80 port and target group as created.

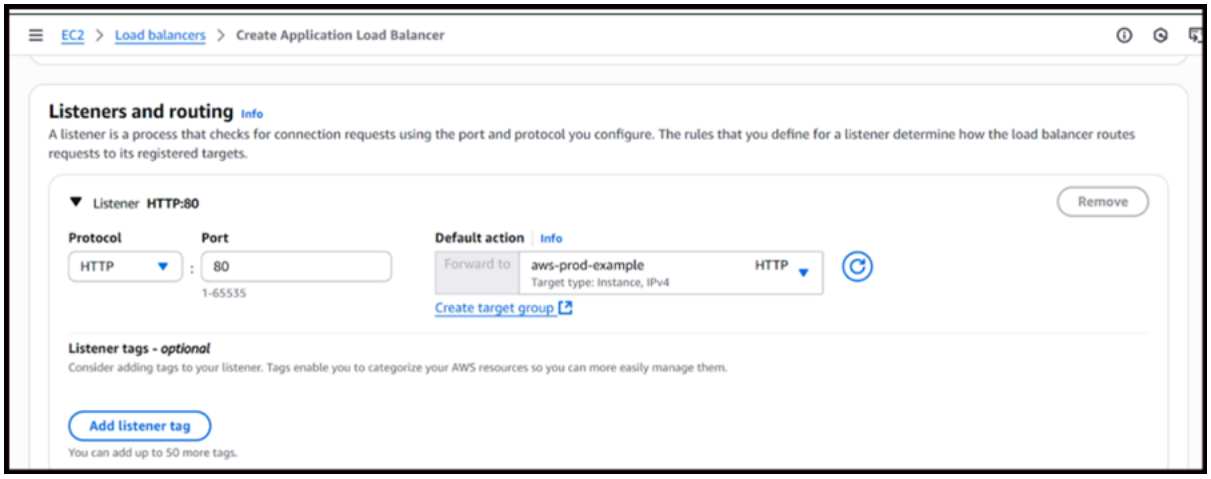

When you try to access the application at port, you won't allow it because the load balancer is allowing port 80 access.

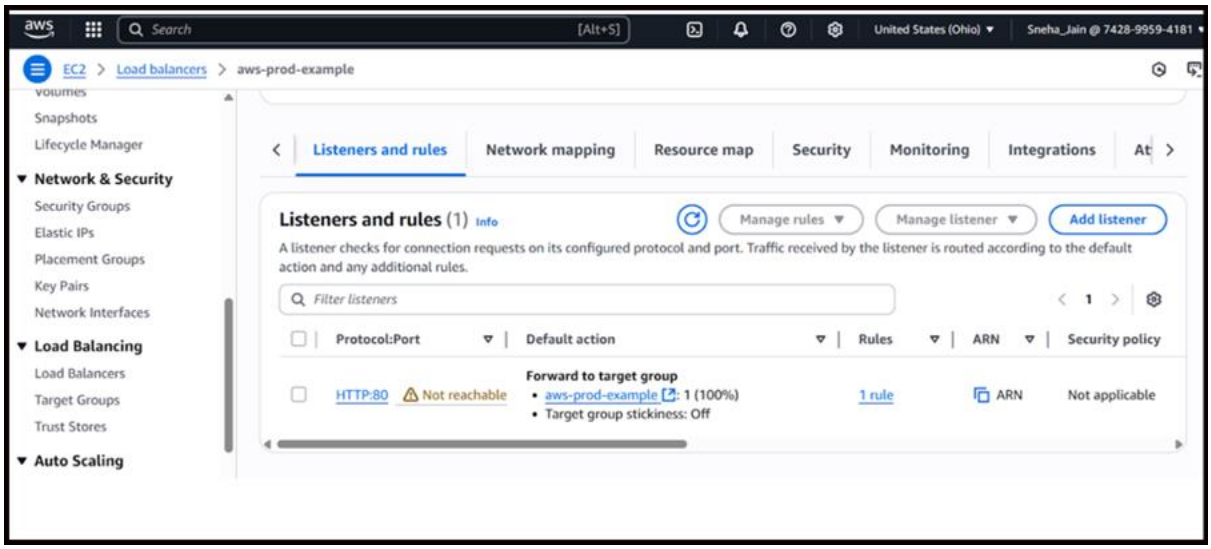

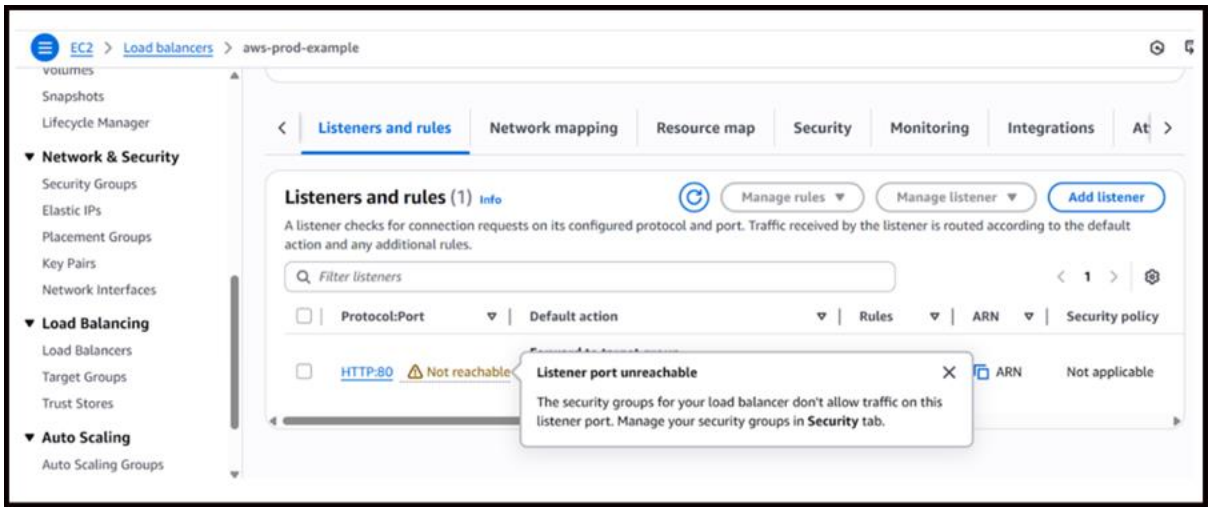

Go to the security group and allow the traffic for port 80.

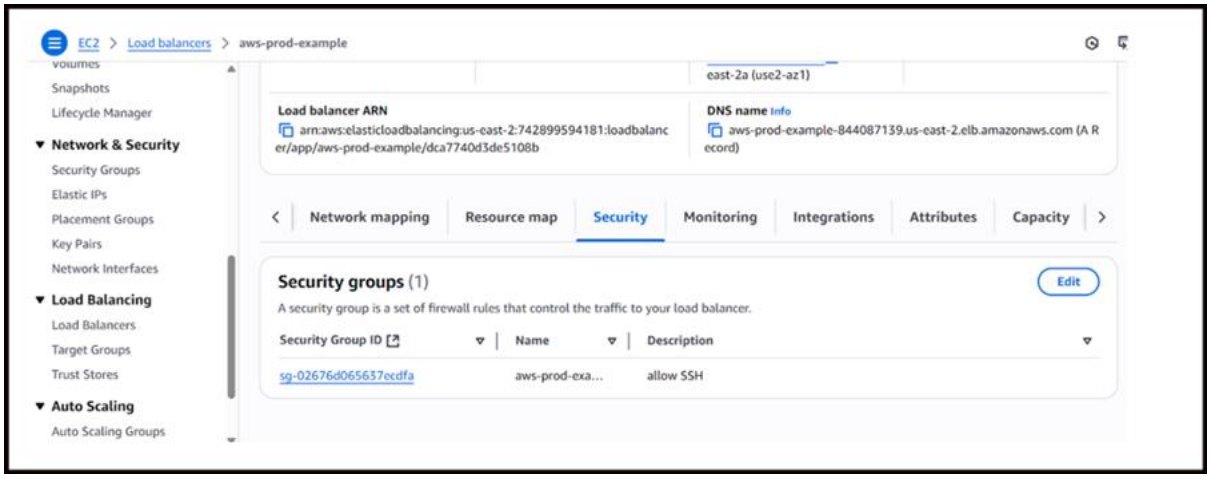

Go to inbound traffic rules and add a port 80 configuration.

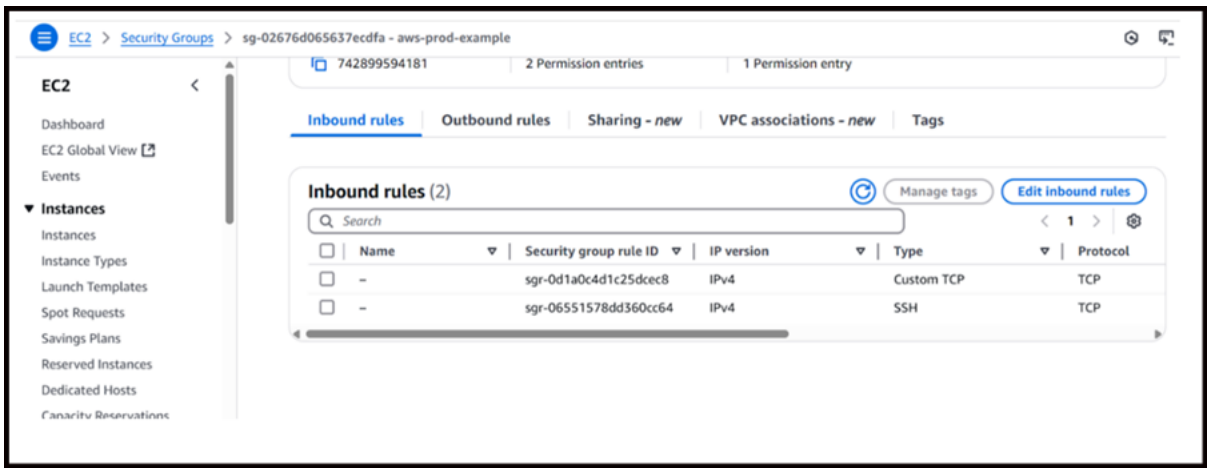

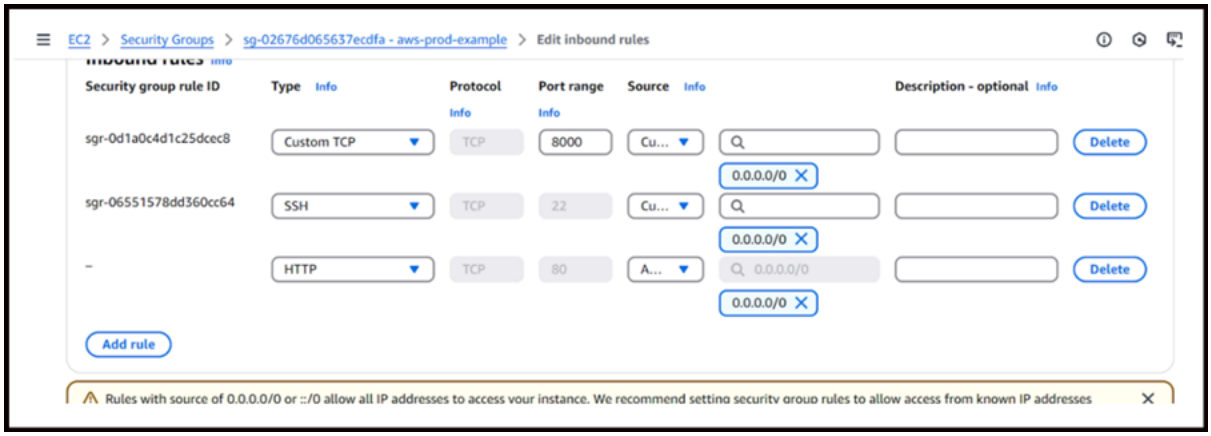

save this rule. and after adding this 80 port in group the error in load balancer disappeared.

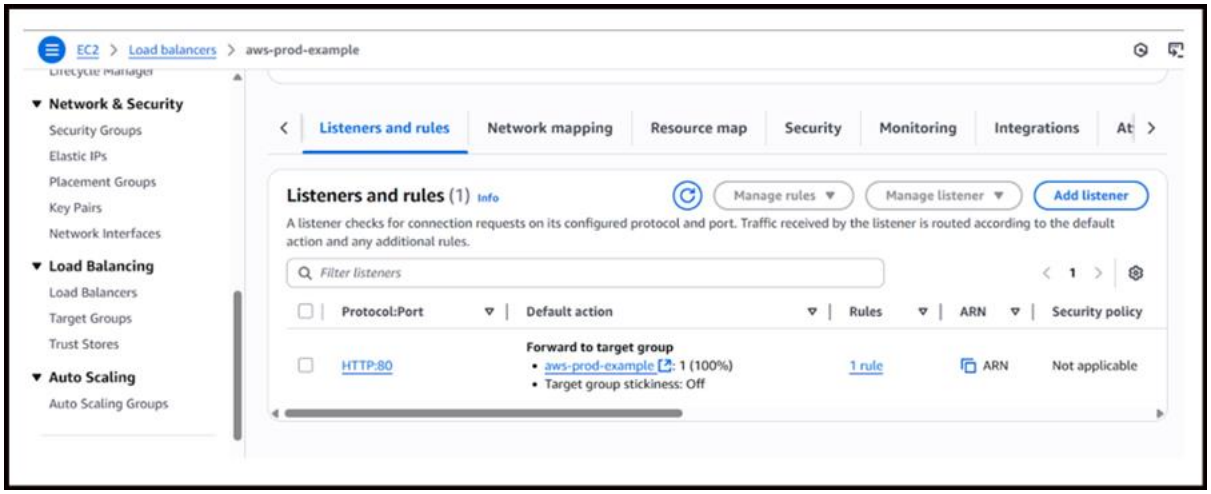

Now access the DNS for the server.

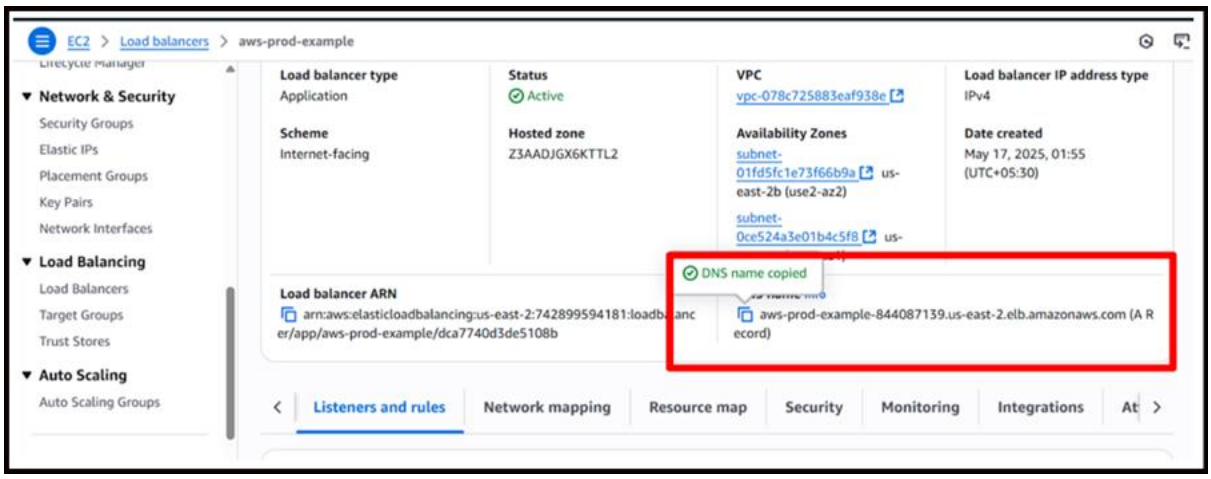

Your application in the private subnet server has deployed now.

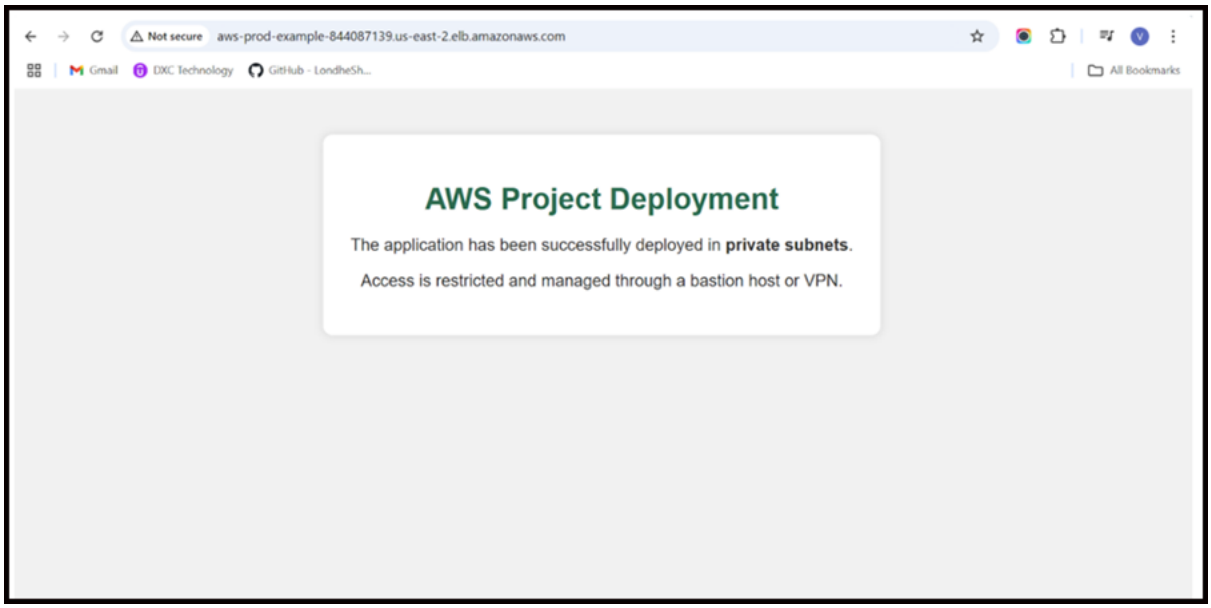

Enjoy:)